Journal de Projet - Déploiement IA

Date : 25 Novembre 2025
Auteur : Michel DONGMO

Contexte

Aider les villes a predire la qualité de l'air en se basant sur les anciennes données

Objectif

Préparer et containeriser l'API de prédiction pour un déploiement public et vérifier le fonctionnement local.

Choix Techniques

- Modèle : Random Forest (choisi pour son F1-score de 0.8284165091158822).
- API : FastAPI choisi pour sa rapidité et sa documentation automatique.
- Conteneurisation : Docker utilisé pour garantir la portabilité.

Étapes réalisées

1. Export du modèle depuis le notebook.
2. Création de l'API FastAPI (`app.py`) exposant `/health` et `/predict`.
3. Ajout des dépendances dans `requirements.txt`.
4. Ajout du modèle `model/best_model.pkl`.
5. Rédaction du Dockerfile et .dockerignore.
6. Tests locaux via `uvicorn` et `curl`.
7. Construction et exécution de l'image Docker, tests via `docker ps` et `curl`.
8. Réalisation des captures d'écran et génération du PDF.

Choix techniques

- FastAPI / Uvicorn pour performance et docs automatiques
- Pickle / joblib pour sérialisation; ONNX possible pour portabilité
- Image basée sur `python:3.11-slim`

Problèmes rencontrés

- Erreur lors du décodage des variables encodées par fréquence
- Erreur de désérialisation XGBoost -> solution: re-sérialiser avec joblib ou exporter en ONNX.
- Durée d'optimisation très longue à cause des variétés des hyperparamètres

Tests effectués

- GET /health -> ok
- POST /predict -> pas ok

Conclusion & recommandations

Convertir en ONNX pour portabilité
Ajouter tests unitaires, CI/CD)

Illustrations

Script Bash pour la création de l'environnement et autres commandes

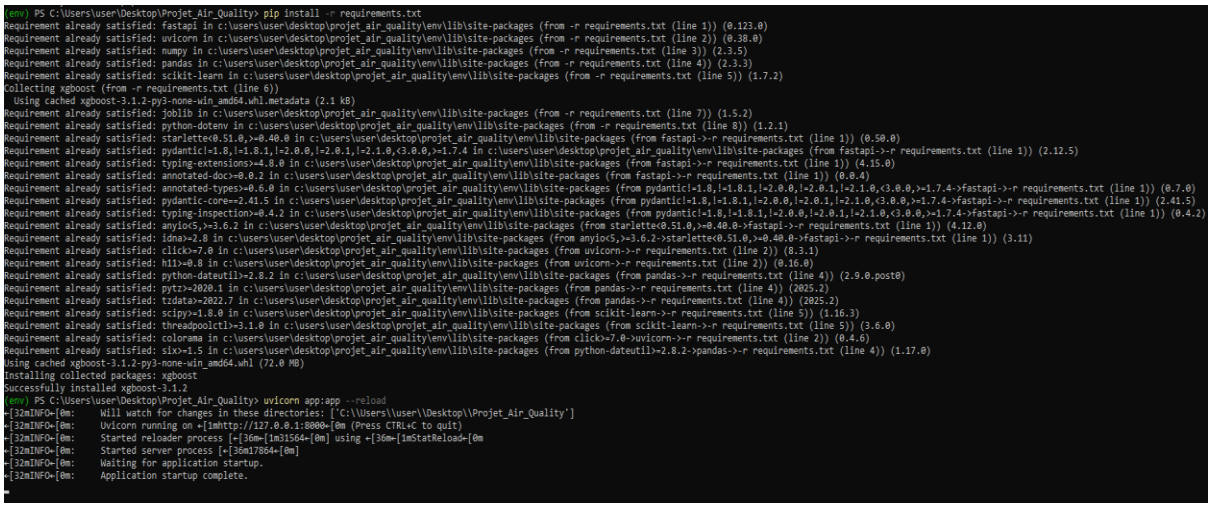

Arborescence du projet avec environnement

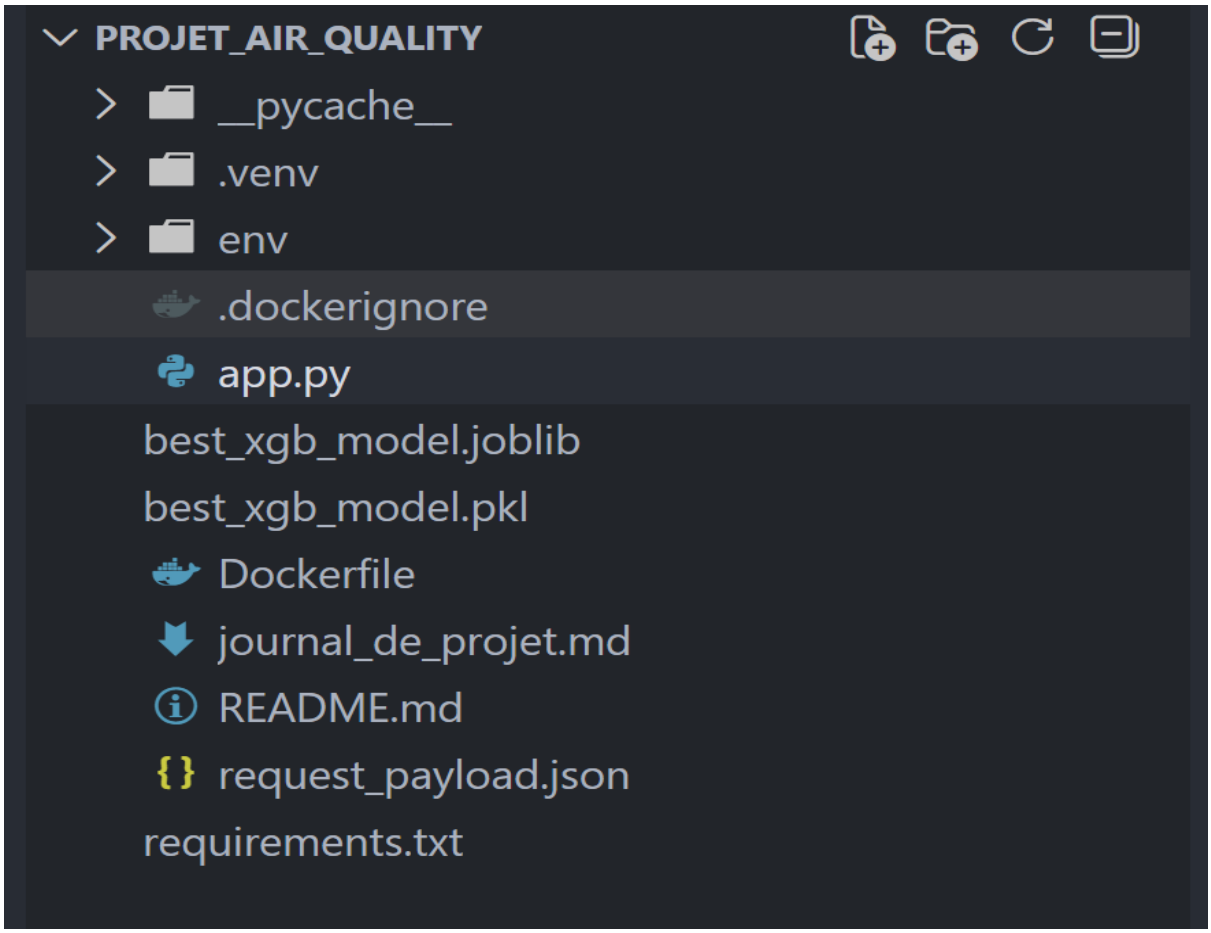

Interface de l'API sur un navigateur web

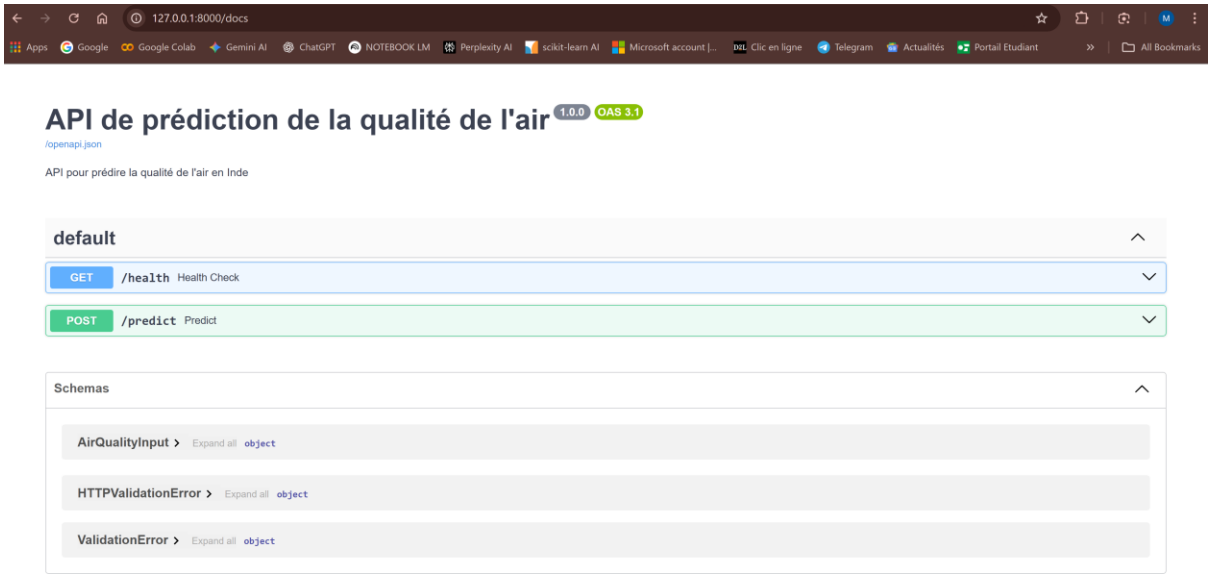

Lancement d'un script json sur try out

Réponse du try out

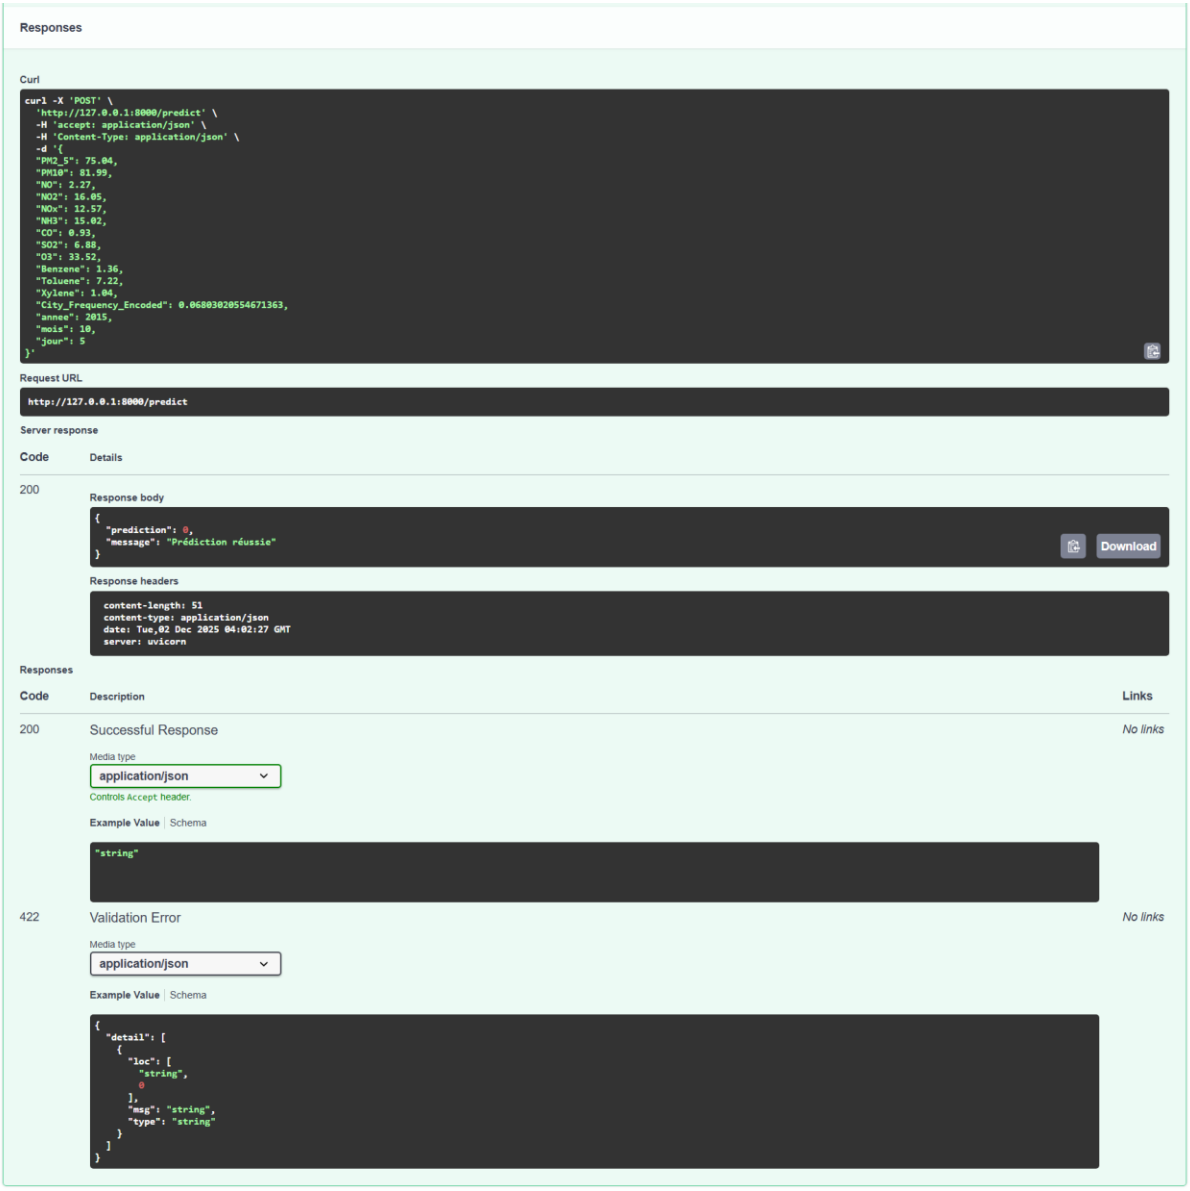

Reponse de l'API

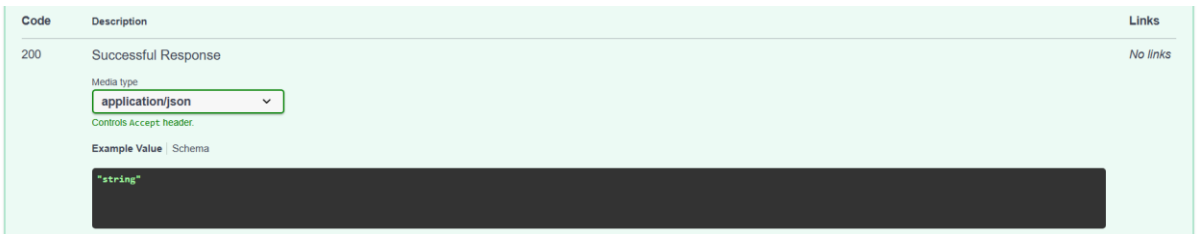

Arrêt du serveur avec la commande << Ctrl + C >> après avoir testé l'API

```
(env) PS C:\Users\user\Desktop\Projet_Air_Quality> uvicorn app:app --reload
←[32mINFO←[0m:     Will watch for changes in these directories: ['C:\\Users\\user\\Desktop\\Projet_Air_Quality']
←[32mINFO←[0m:     Uvicorn running on ←[1mhttp://127.0.0.1:8000←[0m (Press CTRL+C to quit)
←[32mINFO←[0m:     Started reloader process [←[36m←[1m31564←[0m] using ←[36m←[1mStatReload←[0m
←[32mINFO←[0m:     Started server process [←[36m17864←[0m]
←[32mINFO←[0m:     Waiting for application startup.
←[32mINFO←[0m:     Application startup complete.
←[32mINFO←[0m:     127.0.0.1:57686 - "←[1mGET /docs HTTP/1.1←[0m" ←[32m200 OK←[0m
←[33mWARNING←[0m:  StatReload detected changes in 'app.py'. Reloading...
←[32mINFO←[0m:     127.0.0.1:51262 - "←[1mGET / HTTP/1.1←[0m" ←[31m404 Not Found←[0m
←[32mINFO←[0m:     127.0.0.1:50061 - "←[1mGET /docs HTTP/1.1←[0m" ←[32m200 OK←[0m
←[32mINFO←[0m:     127.0.0.1:51034 - "←[1mGET /docs HTTP/1.1←[0m" ←[32m200 OK←[0m
 ←[32mINFO←[0m:     Shutting down
←[32mINFO←[0m:     Waiting for application shutdown.
←[32mINFO←[0m:     Application shutdown complete.
←[32mINFO←[0m:     Finished server process [←[36m17864←[0m]
←[32mINFO←[0m:     Started server process [←[36m10204←[0m]
←[32mINFO←[0m:     Waiting for application startup.
←[32mINFO←[0m:     Application startup complete.
←[32mINFO←[0m:     127.0.0.1:50276 - "←[1mGET /docs HTTP/1.1←[0m" ←[32m200 OK←[0m
←[32mINFO←[0m:     127.0.0.1:50276 - "←[1mGET /openapi.json HTTP/1.1←[0m" ←[32m200 OK←[0m
←[32mINFO←[0m:     127.0.0.1:57677 - "←[1mPOST /predict HTTP/1.1←[0m" ←[32m200 OK←[0m
←[32mINFO←[0m:     Shutting down
←[32mINFO←[0m:     Waiting for application shutdown.
←[32mINFO←[0m:     Application shutdown complete.
←[32mINFO←[0m:     Finished server process [←[36m10204←[0m]
←[32mINFO←[0m:     Stopping reloader process [←[36m←[1m31564←[0m]
(env) PS C:\Users\user\Desktop\Projet_Air_Quality> _
```

Test 2 : Construction et Exécution avec Docker

On commence par lancer l'application docker

1. **Compiler (Build) :**

   docker build -t air-quality-api .

```
(env) PS C:\Users\user\Desktop\Projet_Air_Quality> docker build -t air-quality-api .
[+] Building 130.5s (10/10) FINISHED                                docker:desktop-linux
 => [internal] load build definition from Dockerfile                                0.1s
 => => transferring dockerfile: 557B                                                0.0s
 => [internal] load metadata for docker.io/library/python:3.9-slim                  1.2s
 => [internal] load .dockerignore                                                   0.1s
 => => transferring context: 99B                                                    0.0s
 => [1/5] FROM docker.io/library/python:3.9-slim@sha256:2d97f6910b16bd338d3060f261f53  5.8s
 => => resolve docker.io/library/python:3.9-slim@sha256:2d97f6910b16bd338d3060f261f53  0.1s
 => => sha256:ea56f685404adf81680322f152d2cfec62115b30dda481c2c450078315b 251B / 251B  0.1s
 => => sha256:fc74430849022d13b0d44b8969a953f842f59c6e9d1a0c2c83d71 13.88MB / 13.88MB  1.3s
 => => sha256:b3ec39b36ae8c03a3e09854de4ec4aa08381dfed84a9daa075048c2 1.29MB / 1.29MB  0.8s
 => => sha256:38513bd7256313495cdd83b3b0915a633cfa475dc2a07072ab2c8 29.78MB / 29.78MB  3.4s
 => => extracting sha256:38513bd7256313495cdd83b3b0915a633cfa475dc2a07072ab2c8d191020  1.1s
 => => extracting sha256:b3ec39b36ae8c03a3e09854de4ec4aa08381dfed84a9daa075048c2e3df3  0.2s
 => => extracting sha256:fc74430849022d13b0d44b8969a953f842f59c6e9d1a0c2c83d710affa28  0.7s
 => => extracting sha256:ea56f685404adf81680322f152d2cfec62115b30dda481c2c450078315be  0.0s
 => [internal] load build context                                                   2.1s
 => => transferring context: 14.44MB                                                2.0s
 => [2/5] WORKDIR /app                                                              0.2s
 => [3/5] COPY requirements.txt .                                                   0.1s
 => [4/5] RUN pip install --no-cache-dir -r requirements.txt                       61.5s
 => [5/5] COPY . .                                                                  0.2s
 => exporting to image                                                             60.8s
 => => exporting layers                                                            50.5s
 => => exporting manifest sha256:240d74f24a1a09e5805209b81498f554cc73429c12a3f599b4d4  0.0s
 => => exporting config sha256:1214b2e48573b8a3f695cedede32de015a0e56cd74de74695486da  0.0s
 => => exporting attestation manifest sha256:cd660ea09eaaaea98bb0cb4975b55a8a1c9690c5  0.1s
 => => exporting manifest list sha256:fa0f0f88d12eade535566d8d8cb963118e52f007a73ae96  0.0s
 => => naming to docker.io/library/air-quality-api:latest                           0.0s
 => => unpacking to docker.io/library/air-quality-api:latest                       10.0s

View build details: docker-desktop://dashboard/build/desktop-linux/desktop-linux/2zdkpcp1uqkocuwpoi3idry2u
```

2. Exécuter (Run) :
   docker run -p 8000:8000 –name mon-container-ia air-quality-api

```
(env) PS C:\Users\user\Desktop\Projet_Air_Quality> docker run -p 8000:8000 --na
me mon-container-ia air-quality-api
/usr/local/lib/python3.9/site-packages/xgboost/core.py:158: UserWarning: [04:42
:57] WARNING: /workspace/src/collective/../data/../common/error_msg.h:80: If yo
u are loading a serialized model (like pickle in Python, RDS in R) or
configuration generated by an older version of XGBoost, please export the model
 by calling
`Booster.save_model` from that version first, then load it back in current vers
ion. See:

    https://xgboost.readthedocs.io/en/stable/tutorials/saving_model.html

for more details about differences between saving model and serializing.

  warnings.warn(smsg, UserWarning)
INFO:     Started server process [1]
INFO:     Waiting for application startup.
INFO:     Application startup complete.
INFO:     Uvicorn running on http://0.0.0.0:8000 (Press CTRL+C to quit)
```

3. Vérifier que le moteur docker tourne en exécutant la commande suivante dans un deuxième terminale

   docker ps

```
PS C:\Users\user\Desktop\Projet_Air_Quality> docker ps
CONTAINER ID   IMAGE             COMMAND                  CREATED         STATUS         PORTS
               NAMES
7f4c4b4a3f1a   air-quality-api   "uvicorn app:app --h…"   4 minutes ago   Up 4 minutes   0.0.0.0:8000->8000/tcp, [::]:80
00->8000/tcp   mon-container-ia
PS C:\Users\user\Desktop\Projet_Air_Quality>
```

Affichage dans le terminal lorsque le lien http://127.0.0.1:8000/docs est lancé avec docker

Try out de GET /health

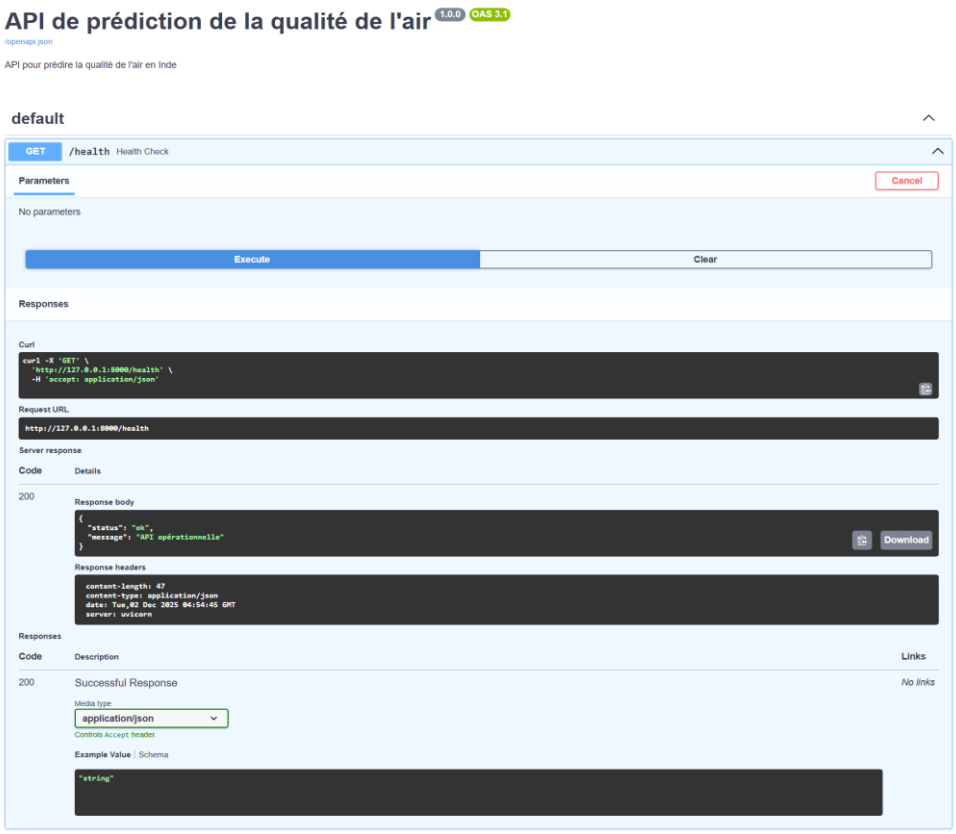

Try out de POST Predict